



Red Hat Ansible Automation Platform 2.6

Automation execution API overview

Developer overview for the automation controller API

Red Hat Ansible Automation Platform 2.6 Automation execution API overview

Developer overview for the automation controller API

Legal Notice

Copyright © Red Hat.

Except as otherwise noted below, the text of and illustrations in this documentation are licensed by Red Hat under the Creative Commons Attribution–Share Alike 3.0 Unported license . If you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, the Red Hat logo, JBoss, Hibernate, and RHCE are trademarks or registered trademarks of Red Hat, LLC. or its subsidiaries in the United States and other countries.

Linux[®] is the registered trademark of Linus Torvalds in the United States and other countries.

XFS is a trademark or registered trademark of Hewlett Packard Enterprise Development LP or its subsidiaries in the United States and other countries.

The OpenStack[®] Word Mark and OpenStack logo are trademarks or registered trademarks of the Linux Foundation, used under license.

All other trademarks are the property of their respective owners.

Abstract

Explore the automation controller API Overview for streamlined automation solutions, empowering developers and administrators with efficient infrastructure management.

Table of Contents

PREFACE	3
PROVIDING FEEDBACK ON RED HAT DOCUMENTATION	4
CHAPTER 1. AVAILABLE TOOLS WITH THE API	5
CHAPTER 2. BROWSING WITH THE API	6
CHAPTER 3. USE CONVENTIONS IN THE API	7
CHAPTER 4. SORTING IN THE API	8
CHAPTER 5. USING THE SEARCH QUERY STRING PARAMETER	9
CHAPTER 6. FILTERING IN THE API	10
6.1. ADVANCED QUERIES IN THE API	10
6.2. FIELD LOOKUPS	11
CHAPTER 7. USING PAGINATION IN THE API	13
CHAPTER 8. ACCESS RESOURCES	14
8.1. CONFIGURATION SETTINGS	14
8.2. IDENTIFIER FORMAT PROTOCOL	15
CHAPTER 9. READ-ONLY FIELDS	17
CHAPTER 10. AUTHENTICATING IN THE API	18
10.1. USING SESSION AUTHENTICATION	18
10.2. BASIC AUTHENTICATION	20
10.2.1. Disabling basic authentication	20
10.3. OAUTH 2 TOKEN AUTHENTICATION	20
10.3.1. Enabling external users to create OAuth 2 tokens	21
10.4. SINGLE SIGN-ON AUTHENTICATION	22

PREFACE

Thank you for your interest in Red Hat Ansible Automation Platform. Ansible Automation Platform helps teams manage complex multitiered deployments by adding control, knowledge, and delegation to Ansible-powered environments.

The automation controller API Overview focuses on helping you understand the automation controller API.

PROVIDING FEEDBACK ON RED HAT DOCUMENTATION

If you have a suggestion to improve this documentation, or find an error, you can contact technical support at <https://access.redhat.com> to open a request.

CHAPTER 1. AVAILABLE TOOLS WITH THE API

Representational State Transfer (REST) relies on a stateless, client/server, and cacheable communications protocol, usually the HTTP protocol.

You might find it helpful to see which API calls the user interface makes in sequence. To do this, you can use the UI from Firebug or Chrome with developer plugins.

Another option is to use [Charles Proxy](#). This offers a visualizer that you might find helpful. While it is commercial software, it can insert itself as an operating system X proxy and intercept both requests from web browsers, curl and other API consumers.

Further options include the following:

- [Fiddler](#)
- [mitmproxy](#)
- [Paros](#)

CHAPTER 2. BROWSING WITH THE API

REST APIs give access to resources (data entities) through URI paths.

Learn how to access and use the automation controller REST API by using a web browser. Learn how to find the recommended current version (v2), view documentation for endpoints, and use GET, PUT, and POST methods with JSON formatting.

Procedure

1. Go to the automation controller REST API in a web browser at:
`https://<gateway server name>/api/controller/v2`
2. Click the "**v2**" link next to "**current versions**" or "**available versions**". Automation controller supports version 2 of the API.
3. Perform a **GET** with just the `/api/` endpoint to get the **current_version**, which is the recommended version.
4. Click the  icon on the navigation menu, for documentation on the access methods for that particular API endpoint and what data is returned when using those methods.
5. Use the **PUT** and **POST** verbs on the specific API pages by formatting JSON in the various text fields.



NOTE

You can also view changed settings from factory defaults at `/api/gateway/v1/settings/` endpoint. It reflects changes you made in the API browser, not changed settings that come from static settings files.


CHAPTER 3. USE CONVENTIONS IN THE API

The API is versioned for compatibility reasons. You can see what API versions are available by querying `/api/`.

You might have to specify the content or type on **POST** or **PUT** requests:

- **PUT**: Update a specific resource (by an identifier) or a collection of resources. You can also use **PUT** to create a specific resource if you know the resource identifier before-hand.
- **POST**: Create a new resource. Also acts as a catch-all verb for operations that do not fit into the other categories.

All URLs not ending with "/" receive a 301 redirect.



NOTE

The formatting of **extra_vars** attached to Job Template records is preserved. YAML is returned as YAML with formatting and comments preserved, and JSON is returned as JSON.

CHAPTER 4. SORTING IN THE API

To give you examples that are easy to follow, we use the following URL throughout this guide:

```
https://<gateway server name>/api/controller/v2/groups/
```

Learn how to use the `order_by` query parameter in a GET request to sort the returned `{{model_verbose_name_plural}}` by one or more fields, including how to specify reverse order by using a leading dash (-).

Procedure

- To specify that `{{ model_verbose_name_plural }}` are returned in a particular order, use the **order_by** query string parameter on the **GET** request:

```
https://<gateway server name>/api/controller/v2/model_verbose_name_plural?order_by={{order_field }}
```

- Prefix the field name with a dash (-) to sort in reverse:

```
https://<gateway server name>/api/controller/v2/model_verbose_name_plural?order_by=-{{ order_field }}
```

- You can specify the sorting fields by separating the field names with a comma (,):

```
https://<gateway server name>/api/controller/v2/model_verbose_name_plural?order_by={{ order_field }},some_other_field
```

CHAPTER 5. USING THE SEARCH QUERY STRING PARAMETER

The search query string parameter provides a simple yet powerful way to filter and find specific resources within your Ansible Automation Platform API. Use a non-case-sensitive search to locate relevant data across designated text fields in a model and extend your search across related fields to find exactly what you need.

Learn how to use the search query parameter to perform a general, non-case-sensitive search across all text fields of an API model and how to use the **related__search** parameter to search across related fields.

Procedure

- Use the search query string parameter to perform a non-case-sensitive search within all designated text fields of a model:

```
https://<gateway server name>/api/controller/v2/model_verbose_name?search=findme
```

- To search across related fields, use the following:

```
https://<gateway_server_name>/api/controller/v2/model_verbose_name?  
related__search=findme
```

CHAPTER 6. FILTERING IN THE API

The system recognizes a collection as a "queryset". You can filter this by using various operators.

Learn how to use query parameters within the REST API URL to filter resources based on criteria such as exact matches, partial string containment (**__contains**), integer casting (**__int**), and many fields simultaneously.

Procedure

- To find groups that contain the name "car", use the following:

```
http://<gateway server name>/api/controller/v2/groups/?name__contains=car
```

- To find an exact match, use the following:

```
http://<gateway server name>/api/controller/v2/groups/?name=car
```

- If a resource is of an integer type, you must add **__int** to the end to cast your string input value to an integer, such as the following:

```
http://<gateway server name>/api/controller/v2/arbitrary_resource/?x__int=5
```

- You can query related resources with the following:

```
http://<gateway server name>/api/gateway/v1/users/?first_name__icontains=kim
```

This returns all users with names that include the string "kim" in them.

- You can also filter against many fields at once:

```
http://<gateway server name>/api/controller/v2/groups/?
name__icontains=test&has_active_failures=false
```

This finds all groups containing the name "test" that have no active failures.

You can also watch the API as the UI is being used to see how it is filtering on various criteria.

Additional resources

- [QuerySet API reference](#)

6.1. ADVANCED QUERIES IN THE API

You can use additional query string parameters used to filter the list of results returned to those matching a given value. You can only use fields and relations that exist in the database for filtering.

Ensure that any special characters in the specified value are URL-encoded. For example:

```
?field=value%20xyz
```

Fields can also span relations, only for fields and relationships defined in the database:

■

```
?other__field=value
```

To exclude results matching certain criteria, prefix the field parameter with **not__**:

```
?not__field=value
```

By default, all query string filters are AND'ed together, so only the results matching all filters are returned. To combine results matching any one of multiple criteria, prefix each query string parameter with **or__**:

```
?or__field=value&or__field=othervalue  
?or__not__field=value&or__field=othervalue
```

The default AND filtering applies all filters simultaneously to each related object being filtered across database relationships. The chain filter instead applies filters separately for each related object. To use this, prefix the query string parameter with **chain__**:

```
?chain__related__field=value&chain__related__field2=othervalue  
?chain__not__related__field=value&chain__related__field2=othervalue
```

If you write the first query as **?relatedfield=value&relatedfield2=othervalue**, it returns only the primary objects where the same related object satisfied both conditions. As written by using the chain filter, it would return the intersection of primary objects matching each condition.

6.2. FIELD LOOKUPS

Field lookups allow you to filter API results based on specific criteria. You can use field lookups for more advanced queries, by appending the lookup to the field name:

```
?field__lookup=value
```

The following field lookups are supported:

- **exact**: Exact match (default lookup if not specified, see the following note for more information).
- **ixact**: Case-insensitive version of exact.
- **contains**: Field contains value.
- **icontains**: Case-insensitive version of contains.
- **startswith**: Field starts with value.
- **istartswith**: Case-insensitive version of startswith.
- **endswith**: Field ends with value.
- **iendswith**: Case-insensitive version of endswith.
- **regex**: Field matches the given regular expression.
- **iregex**: Case-insensitive version of regular expression.

- `gt`: Greater than comparison.
- `gte`: Greater than or equal to comparison.
- `lt`: Less than comparison.
- `lte`: Less than or equal to comparison.
- `isnull`: Check whether the given field or related object is null; expects a boolean value.
- `in`: Check whether the given field's value is present in the list provided; expects a list of items.
- You can specify boolean values as **True** or **1** for true, **False** or **0** for false (both case-insensitive).

For example, `?created__gte=2023-01-01` provides a list of items created after 1/1/2023.

You can specify null values as **None** or **Null** (both case-insensitive), though we recommend using the **isnull** lookup to explicitly check for null values.

You can specify lists (for the **in** lookup) as a comma-separated list of values. Filtering based on the requesting user's level of access by query string parameter:

- **role_level**: Level of role to filter on, such as **admin_role**



NOTE

Earlier releases of Ansible Automation Platform returned queries with **_exact** results by default. As a workaround, set the **limit** to **?limit_exact** for the default filter. For example, `/api/controller/v2/jobs/?limit_exact=example.domain.com` results in:

```
{
  "count": 1,
  "next": null,
  "previous": null,
  "results": [
    ...
  ]
}
```

CHAPTER 7. USING PAGINATION IN THE API

The API paginates responses for collections. This means that while a collection might contain tens or hundreds of thousands of objects, in each web request, only a limited number of results are returned for API performance reasons.


When you receive the result for a collection, something similar to the following is displayed:

```
{'count': 25, 'next': 'http://testserver/api/controller/v2/some_resource?page=2', 'previous': None,
 'results': [ ... ] }
```

Procedure

1. Request the page given by the "next" sequential URL, to get to the next page.
2. Use the **page_size=XX** query string parameter to change the number of results returned for each request.
 - The **page_size** has a default maximum limit of 200, which is enforced when a user tries a value beyond it, for example, **?page_size=1000**. However, you can change this limit by setting the value in **/etc/tower/conf.d/<some file>.py** to something higher. For example, **MAX_PAGE_SIZE=1000**.
3. Use the page query string parameter to retrieve a particular page of results:

```
http://<gateway server name>/api/controller/v2/model_verbose_name?
page_size=100&page=2
```



NOTE

The preceding and following links returned with the results set these query string parameters automatically.

Do not request page sizes greater than 200.

The user interface uses smaller values to avoid scrolling.

CHAPTER 8. ACCESS RESOURCES

Automation controller uses a primary key to access individual resource objects. You can access automation controller resources by using resource-specific, human-readable identifiers through the named URL feature.

The following example shows the named URL path where you can access a resource object without an auxiliary query string:

```
/api/controller/v2/hosts/host_name++inv_name++org_name/
```

8.1. CONFIGURATION SETTINGS

There are two named-URL-related configuration settings available under `/api/controller/v2/settings/named-url/`: **NAMED_URL_FORMATS** and **NAMED_URL_GRAPH_NODES**.

NAMED_URL_FORMATS is a read only key-value pair list of all available named URL identifier formats. The following shows an example **NAMED_URL_FORMATS**:

```
"NAMED_URL_FORMATS": {
  "organizations": "<name>",
  "teams": "<name>++<organization.name>",
  "credential_types": "<name>+<kind>",
  "credentials": "<name>++<credential_type.name>+<credential_type.kind>++<organization.name>",
  "notification_templates": "<name>++<organization.name>",
  "job_templates": "<name>++<organization.name>",
  "projects": "<name>++<organization.name>",
  "inventories": "<name>++<organization.name>",
  "hosts": "<name>++<inventory.name>++<organization.name>",
  "groups": "<name>++<inventory.name>++<organization.name>",
  "inventory_sources": "<name>++<inventory.name>++<organization.name>",
  "inventory_scripts": "<name>++<organization.name>",
  "instance_groups": "<name>",
  "labels": "<name>++<organization.name>",
  "workflow_job_templates": "<name>++<organization.name>",
  "workflow_job_template_nodes": "<identifier>++<workflow_job_template.name>++<organization.name>",
  "applications": "<name>++<organization.name>",
  "users": "<username>",
  "instances": "<hostname>"
}
```

For each item in **NAMED_URL_FORMATS**, the key is the API name of the resource to have named URL. The value is a string indicating how to form a human-readable unique identifier for that resource. **NAMED_URL_FORMATS** only lists the resources that can have named URL, any resource not listed there has no named URL. If a resource can have named URL, its objects must have a **named_url** field that represents the object-specific named URL. That field is only visible under detail view, not list view. You can access specified resource objects by using accurately generated named URL. This the object and its related URLs. For example, if `/api/controller/v2/res_name/obj_slug/` is valid, `/api/controller/v2/res_name/obj_slug/related_res_name/` is also valid.

NAMED_URL_FORMATS are instructive enough to compose human-readable unique identifiers and named URLs themselves. For ease-of-use, every object of a resource that can have named URL has a

related field **named_url** that displays that object's named URL. You can copy and paste that field for your own custom use. For more information, see the help text of the API browser if a resource object has named URL.

You can manually decide the named URL label, for example with ID 5. To compose a named URL for this specific resource object by using **NAMED_URL_FORMATS**, first look up the labels field of **NAMED_URL_FORMATS** to get the identifier format **<name>++<organization.name>**:

- The first part of the URL format is **<name>**, which indicates that you can find the label resource detail in **/api/controller/v2/labels/5/**, and look for **name** field in returned JSON. If you have the **name** field with value **Foo**, then the first part of the unique identifier is **Foo**.
- The second part of the format is double plus sign **++**. That is the delimiter that separates different parts of a unique identifier. Append them to the unique identifier to get **Foo++**.
- The third part of the format is **<organization.name>**, which indicates that the field is not in the current label object under investigation, but in an organization that the label object points to. As the format indicates, look up the organization in the related field of the current returned JSON. That field might not exist. If it exists, follow the URL given in that field, for example, **/api/gateway/v1/organizations/3/**, to get the details of the specific organization, extract its name field, for example, "Default", and append it to the current unique identifier. Since **<organizations.name>** is the last part of the format, it generates the following named URL: **/api/controller/v2/labels/Foo++Default/**.
In the case where an organization does not exist in the related field of the label object detail, append an empty string instead. This does not alter the current identifier. Therefore, **Foo++** becomes the final unique identifier and the resulting generated named URL becomes **/api/controller/v2/labels/Foo++/**.

An important aspect of generating a unique identifier for named URL has to do with reserved characters. As the identifier is part of a URL, the following reserved characters by URL standard are encoded by percentage symbols: **;/?:@=&[]**. For example, if an organization is named **;/?:@=&[]**, its unique identifier should be **%3B%2F%3F%3A%40%3D%26%5B%5D**. Another special reserved character is **+**, which is not reserved by URL standard but used by named URL to link different parts of an identifier. It is encoded by **[+]**. For example, if an organization is named **[+]**, its unique identifier is **%5B[+]%5D**, where original **[** and **]** are percent encoded and **+** is converted to **[+]**.

Although you cannot manually change **NAMED_URL_FORMATS**, modifications do occur automatically and expand over time, reflecting underlying resource modification and expansion. Consult the **NAMED_URL_FORMATS** on the same cluster where you want to use the named URL feature.

NAMED_URL_GRAPH_NODES is another read-only list of key-value pairs that exposes the internal graph data structure used to manage named URLs. This is not intended to be human-readable but must be used for programmatically generating named URLs. An example script for generating named URL given the primary key of arbitrary resource objects that can have a named URL, using info provided by **NAMED_URL_GRAPH_NODES**, can be found in [GitHub](#).

8.2. IDENTIFIER FORMAT PROTOCOL

automation controller uses a consistent protocol to generate human-readable unique identifiers for resources that can be addressed by name in the API and the web interface.

Resources are identifiable by their unique keys, which are tuples of resource fields. Every resource is guaranteed to have its primary key number alone as a unique key, but there might be many other unique keys. A resource can generate an identifier format and, therefore, have a named URL if it has at least

one unique key that satisfies the following rules:

1. The key must contain only fields that are either the **name** field, or text fields with a finite number of possible choices (such as credential type resource's **kind** field).
2. The only permitted exceptional field that breaks the preceding rule is a many-to-one related field relating to a resource other than itself, which is also allowed to have a slug.

If there are resources **Foo** and **Bar**, both **Foo** and **Bar** contain a name field and a choice field that can only have values "yes" or "no". Additionally, resource **Foo** has a many-to-one field (a foreign key) relating to **Bar**, for example **fk**. **Foo** has a unique key tuple (**name**, **choice**, **fk**) and **Bar** has a unique key tuple (**name**, **choice**). **Bar** can have named URL because it satisfies the preceding first rule. **Foo** can also have named URL, even though it breaks the first rule, the extra field breaking rule number one is the **fk** field, which is many-to-one-related to **Bar** and **Bar** can have named URL.

For resources satisfying rule number one, their human-readable unique identifiers are combinations of foreign key fields, delimited by **+**. In specific, resource **Bar** in the preceding example has slug format **<name>+<choice>**.

Note that the field order matters in slug format and the **name** field always comes first if present, followed by the remaining fields arranged in lexicographic order of field name. For example, if **Bar** also has an **a_choice** field satisfying rule one and the unique key becomes (**name**, **choice**, **a_choice**), its slug format becomes **<name>+<a_choice>+<choice>**.

For resources satisfying rule number two, if traced back through the extra foreign key fields, the result is a tree of resources that identify objects of that resource. To generate the identifier format, each resource in the traceback tree generates its own part of the standalone format, using all fields but the foreign keys. Finally, all parts are combined by **++** in the following order:

- Put standalone format as the first identifier part.
- Recursively generate unique identifiers for each resource. The underlying resource is pointing to using a foreign key (a child of a traceback tree node).
- Treat generated unique identifiers as the rest of the identifier components. Sort them in lexicographic order of corresponding foreign keys.
- Combine all components together using **++** to generate the final identifier format.

When generating an identifier format for resource **Foo**, automation controller generates the standalone formats, **<name>+<choice>** for **Foo** and **<fk.name>+<fk.choice>** for **Bar**, then combines them together to be **<name>+<choice>++<fk.name>+<fk.choice>**.

When generating identifiers according to the given identifier format, there are cases where a foreign key might point to nowhere. In this case, automation controller substitutes the part of the format corresponding to the resource the foreign key should point to with an empty string. For example, if a **Foo** object has the **name = "alice"**, **choice = "yes"**, but **fk field = None**, its resulting identifier is **alice+yes++**.

CHAPTER 9. READ-ONLY FIELDS

Certain fields in the REST API are marked read-only.

These usually include the URL of a resource, the ID, and occasionally some internal fields. For example, the **'created_by'** attribute of each object indicates which user created the resource, and you cannot edit this. If you post some values and notice that they are not changing, these fields might be read-only.

CHAPTER 10. AUTHENTICATING IN THE API

Automation controller is designed for organizations to control their automation with a visual dashboard for out-of-the box control while providing a REST API to integrate with your other tools on a deeper level.

Automation controller supports several authentication methods to make it easy to embed automation controller into existing tools and processes. This ensures that the right people can access its resources.

You can use the following authentication methods in the API:

- [Session authentication](#)
- [Basic authentication](#)
- [OAuth 2 token authentication](#)
- [Single sign-on authentication](#)

10.1. USING SESSION AUTHENTICATION

You can use session authentication when logging in directly to the automation controller's API or UI to manually create resources, such as inventories, projects, and job templates and start jobs in the browser.

With this method, you can remain logged in for a prolonged period of time, not just for that HTTP request. For example, when browsing the API or UI in a browser such as Chrome or Mozilla Firefox. When you log in, a session cookie is created, this means that you can remain logged in when navigating to different pages within automation controller.

The following image represents the communication that occurs between the client and server in a session:



Use the curl tool to see the activity that occurs when you log in to automation controller.

Procedure

1. Use **GET** to go to the `/api/login/` endpoint to get the **csrftoken** cookie:

```
$ curl -k -c - https://<gateway server name>/api/gateway/v1/login/

$YOUR_AAP_URL FALSE / TRUE 1780539778 csrftoken
GODXonA5LyV1uAs8zvcD2k12DQJC74oB
```

2. Extract the **csrftoken** value from the cookie returned in the first step.
3. **POST** to the `/api/login/` endpoint with username, password, and **X-CSRFToken=<token-value>**:

```
curl -X POST -H 'Content-Type: application/x-www-form-urlencoded' \

--referer https://<gateway server name>/api/gateway/v1/login/ \
```

```
-H 'X-CSRFToken: <token-value>' \
--data 'username=admin&password=$YOUR_ADMIN_PASSWORD' \
--cookie 'csrftoken=GODXonA5LyV1uAs8zvcD2k12DQJC74oB' \
https://<gateway server name>/api/gateway/v1/login/ -k -D - -o /dev/null
```

4. Access and test the APIs that need authentication, for example **/api/controller/v2/settings/all/**:



NOTE

All of this is done by automation controller when you log in to the UI or API in the browser, and you must only use it when authenticating in the browser. For programmatic integration with automation controller, see [OAuth2 token authentication](#).

```
$ curl -X GET -H 'Cookie: <cookieID>,' https://<gateway server
name>/api/controller/v2/settings/all/ -k
```

Verification

The following shows a typical response:

```
Server: nginx
Date: <current date>
Content-Type: text/html; charset=utf-8
Content-Length: 0
Connection: keep-alive
Location: /accounts/profile/
X-API-Session-Cookie-Name: awx_sessionid
Expires: <date>
Cache-Control: max-age=0, no-cache, no-store, must-revalidate, private
Vary: Cookie, Accept-Language, Origin
Session-Timeout: 1800
Content-Language: en
X-API-Total-Time: 0.377s
X-API-Request-Id: 700826696425433fb0c8807cd40c00a0
Access-Control-Expose-Headers: X-API-Request-Id
Set-Cookie: userLoggedIn=true; Path=/
Set-Cookie: current_user=<user cookie data>; Path=/
Set-Cookie: csrftoken=<csrftoken>; Path=/; SameSite=Lax
Set-Cookie: awx_sessionid=<your session id>; expires=<date>; HttpOnly; Max-Age=1800; Path=/;
SameSite=Lax
Strict-Transport-Security: max-age=15768000
```

When a user is successfully authenticated with this method, the server responds with a header called **X-API-Session-Cookie-Name**, indicating the configured name of the session cookie. The default value is **awx_session_id** which you can see later in the **Set-Cookie** headers.

**NOTE**

You can change the session expiration time by specifying it in the **SESSION_COOKIE_AGE** parameter.

10.2. BASIC AUTHENTICATION

Basic authentication is stateless, therefore, you must send the base64-encoded username and password along with each request through the Authorization header. You can use this for API calls from curl requests, python scripts, or individual requests to the API.

We recommend OAuth 2 Token Authentication for accessing the API when at all possible.

The following is an example of Basic authentication with curl:

```
# the --user flag adds this Authorization header for us
curl -X GET --user 'user:password' https://<controller-host>/api/gateway/v1/tokens/ -k -L
```

Additional resources

- [The 'Basic' HTTP Authentication Scheme](#)

10.2.1. Disabling basic authentication

You can disable basic authentication for security purposes.

Procedure

1. From the navigation panel, select **Settings** → **Platform gateway**.
2. Click **Edit platform gateway settings**.
3. Disable the option **Gateway basic auth enabled**.
4. Click **Save platform gateway settings**.

10.3. OAUTH 2 TOKEN AUTHENTICATION

OAuth (Open Authorization) is an open standard for token-based authentication and authorization. OAuth 2 authentication is commonly used when interacting with the automation controller API programmatically.

Similar to Basic authentication, you are given an OAuth 2 token with each API request through the Authorization header. Unlike Basic authentication, OAuth 2 tokens have a configurable timeout and are scorable. Tokens have a configurable expiration time and can be easily revoked for one user or for the entire automation controller system by an administrator if needed. You can do this with the [revoke_oauth2_tokens](#) management command, or by using the API as explained in [Revoke an access token](#).

The different methods for obtaining OAuth2 access tokens in automation controller include the following:

- Personal access tokens (PAT)

- Application token: Password grant type
- Application token: Implicit grant type
- Application token: Authorization Code grant type

A user needs to create an OAuth 2 token in the API or in the **Access Management → OAuth Applications** tab of the platform gateway UI. For more information about creating tokens through the UI, see [Adding tokens](#).

For the purpose of this example, use the PAT method for creating a token in the API. After you create it, you can set the scope.



NOTE

You can configure the expiration time of the token system-wide. For more information, see [Configuring access to external applications with token-based authentication](#).

Token authentication is best used for any programmatic use of the automation controller API, such as Python scripts or tools such as curl.

Curl example

```
curl -u user:password -k -X POST https://<platform-host>/api/gateway/v1/tokens/
```

This call returns JSON data with the following:



You can use the value of the **token** property to perform a **GET** request for an automation controller resource, such as Hosts:

```
curl -k -X GET \
-H "Content-Type: application/json" \
-H "Authorization: Bearer <oauth2-token-value>" \
https://<platform-host>/api/controller/v2/hosts/
```

You can also run a job by making a **POST** to the job template that you want to start:

```
curl -k -X POST \
-H "Authorization: Bearer <oauth2-token-value>" \
-H "Content-Type: application/json" \
--data '{"limit" : "ansible"}' \
https://<platform-host>/api/controller/v2/job_templates/14/launch/
```

Additional resources

- [Configuring access to external applications with token-based authentication](#)

10.3.1. Enabling external users to create OAuth 2 tokens

By default, external users such as those created by single sign-on are not able to generate OAuth tokens for security purposes.

Procedure

1. From the navigation panel, select **Settings → Platform gateway**.
2. Select **Edit platform gateway settings**.
3. Enable the option to **Allow external users to create OAuth2 tokens**

10.4. SINGLE SIGN-ON AUTHENTICATION

Single sign-on (SSO) authentication methods are different from other methods because the authentication of the user happens external to automation controller, such as Google SSO, Microsoft Azure SSO, SAML, or GitHub.

You can configure SSO authentication by using automation controller inside a large organization with a central Identity Provider. Once you have configured an SSO method in automation controller, an option for that SSO is available on the login screen. If you click that option, it redirects you to the Identity Provider, in this case GitHub, where you present your credentials. If the Identity Provider verifies you successfully, automation controller makes a user linked to your GitHub user (if this is your first time logging in with this SSO method), and logs you in.

Additional resources

- [Configuring an authentication type](#)